

— INF4820 —  
Algorithms for AI and NLP  
*Semantic Spaces*

Murhaf Fares & Stephan Oepen

Language Technology Group (LTG)

September 22, 2016



“You shall know a word by the company it keeps!”



- ▶ Alcazar?
- ▶ The alcazar did not become a permanent residence for the royal family until 1905
- ▶ The alcazar was built in the tenth century
- ▶ You can also visit the alcazar while the royal family is there

- ▶ Can a program reuse the same intuition to automatically learn word meaning?
  - ▶ By looking at data of actual language use
  - ▶ and without any prior knowledge
- ▶ How can we represent word meaning in a mathematical model?

## Concepts

- ▶ Distributional semantics
- ▶ Vector spaces
- ▶ Semantic spaces

## AKA the contextual theory of meaning

- *Meaning is use.* (Wittgenstein, 1953)
- *You shall know a word by the company it keeps.* (Firth, 1957)
- *The meaning of entities, and the meaning of grammatical relations among them, is related to the restriction of combinations of these entities relative to other entities.* (Harris, 1968)



- ▶ **The hypothesis:** If two words share similar contexts, we can assume that they have similar meanings.
- ▶ Comparing meaning reduced to comparing contexts,
  - no need for prior knowledge!
- ▶ **Our goal:** to automatically learn word semantics based on this hypothesis.

## A distributional approach to lexical semantics:

- ▶ Given the set of words in our vocabulary  $|V|$
- ▶ Record contexts of words across a large collection of texts (corpus).
- ▶ Each word is represented by a set of contextual **features**.
- ▶ Each feature records some property of the observed contexts.
- ▶ Words that are found to have similar features are expected to also have similar meaning.



- ▶ **The hypothesis:** If two **words** share similar **contexts**, we can assume that they have **similar** meanings.
- ▶ How do we define *word*?
- ▶ How do we define *context*?
- ▶ How do we define *similar*?

# What is a word?



Raw:	“The programmer’s programs had been programmed.”
Tokenized:	the programmer ’s programs had been programmed .
Lemmatized:	the programmer ’s program have be program .
W/ stop-list:	programmer program program
Stemmed:	program program program

- ▶ **Tokenization:** Splitting a text into sentences and words or other units.
- ▶ Different levels of abstraction and morphological normalization:
  - ▶ What to do with case, numbers, punctuation, compounds, ...?
  - ▶ Full-form words vs. lemmas vs. stems ...
- ▶ **Stop-list:** filter out closed-class words or function words.
  - ▶ The idea is that only *content words* provide relevant context.



... Tunisian or French cakes and it is marketed. The **bread** may be cooked such as Kessra or Khmira or Harchaya ...

... Chile, cochayuyo. Laver is used to make laver **bread** in Wales where it is known as "bara lawr"; in ...

... and how everyday events such as a Samurai cutting **bread** with his sword are elevated to something special and ...

... used to make the two main food staples of **bread** and beer. Flax plants, uprooted before they started flowering ...

... for milling grain and a small oven for baking the **bread**. Walls were painted white and could be covered with dyed ...

... of the ancients. The staple diet consisted of **bread** and beer, supplemented with vegetables such as onions and garlic ...

... Prayers were made to the goddess Isis. Moldy **bread**, honey and copper salts were also used to prevent ...

... going souling and the baking of special types of **bread** or cakes. In Tirol, cakes are left for them on the table ...

... under bridges, beg in the streets, and steal loaves of **bread**. If the path be beautiful, let us not question where it ...

... When Jesus the Christ, who is the Word and the **bread** of Life, comes a second time, the righteous will be raised ...



“Rose is a rose is a rose is a rose.” *Gertrude Stein*

Three types and ten tokens.

- ▶ Let's say we're extracting (contextual) features for the target *bread* in:

I bake **bread** for breakfast.

## Context windows

- ▶ Context  $\equiv$  neighborhood of  $\pm n$  words left/right of the focus word.
- ▶ **Features** for  $\pm 1$ : {left:bake, right:for}
- ▶ Some variants: distance weighting, *n*grams.

## Bag-of-Words (BoW)

- ▶ Context  $\equiv$  all co-occurring words, ignoring the linear ordering.
- ▶ **Features**: {I, bake, for, breakfast}
- ▶ Some variants: sentence-level, document-level.

I bake **bread** for breakfast.

## Grammatical context

- ▶ Context  $\equiv$  the grammatical relations to other words.
- ▶ Intuition: When words combine in a construction they often impose semantic constraints on each other:  
*... to {drink | pour | spill} some {milk | water | wine} ...*
- ▶ **Features**: `{dir_obj(bake), prep_for(breakfast)}`
- ▶ Requires deeper linguistic analysis than simple BoW approaches.



- ▶ What do we mean by *similar*?
- ▶ *car, road, gas, service, traffic, driver, license*
- ▶ *car, train, bicycle, truck, vehicle, airplane, bus*
- ▶ Relatedness vs. sameness. Or domain vs. content. Or syntagmatic vs. paradigmatic.
- ▶ Similarity in **domain**: {*car, road, gas, service, traffic, driver, license*}
- ▶ Similarity in **content**: {*car, train, bicycle, truck, vehicle, airplane, bus*}
- ▶ The type of context dictates the type of semantic similarity.
- ▶ Broader definitions of context tend to give clues for *domain-based relatedness*.
- ▶ Fine-grained and linguistically informed contexts give clues for *content-based similarity*.



- ▶ Given the different definitions of 'word', 'context' and 'similarity':
- ▶ How exactly should we represent our words and context features?
- ▶ How exactly can we compare the features of different words?

## A distributional approach to lexical semantics:


- ▶ Record contexts of words across a large collection of texts (corpus).
- ▶ Each word is represented by a set of contextual **features**.
- ▶ Each feature records some property of the observed contexts.
- ▶ Words that are found to have similar features are expected to also have similar meaning.

- ▶ Vector space models first appeared in IR.
- ▶ A general algebraic model for representing data based on a spatial metaphor.
- ▶ Each object is represented as a vector (or point) positioned in a coordinate system.
- ▶ Each coordinate (or **dimension**) of the space corresponds to some descriptive and measurable property (**feature**) of the objects.
- ▶ To measure **similarity** of two objects, we can measure their **geometrical distance** / closeness in the model.
- ▶ Vector representations are foundational to a wide range of ML methods.




- ▶ A vector space is defined by a system of  $n$  dimensions or coordinates where points are represented as **real-valued vectors** in the space  $\mathbb{R}^n$ .
- ▶ The most basic example is 2-dimensional Euclidean plane  $\mathbb{R}^2$ .

$$v_1 = [5, 5], v_2 = [1, 8]$$



- ▶ AKA distributional semantic models or word space models.
- ▶ A semantic space is a vector space model where
  - ▶ **points** represent **words**,
  - ▶ **dimensions** represent **context** of use,
  - ▶ and **distance** in the space represents **semantic similarity**.



Dimensions:  $w1, w2, w3$

$$t1 = [2, 1, 2] \in \mathbb{R}^3$$

$$t2 = [1, 1, 1] \in \mathbb{R}^3$$

- ▶ Each word type  $t_i$  is represented by a vector of real-valued features.
- ▶ Our observed feature vectors must be encoded numerically:
  - ▶ Each context feature is mapped to a dimension  $j \in [1, n]$ .
  - ▶ For a given word, the value of a given feature is its number of **co-occurrences** for the corresponding context across our corpus.
- ▶ Let the set of  **$n$  features** describing the lexical contexts of a **word**  $t_i$  be represented as a **feature vector**  $\vec{x}_i = \langle x_{i1}, \dots, x_{in} \rangle$ .

## Example

- ▶ Given a grammatical context, if we assume that:
- ▶ the  **$i$ th** word is *bread* and
- ▶ the  **$j$ th** feature is OBJ\_OF(bake), then
- ▶  **$x_{ij} = 4$**  would mean that we have observed *bread* to be the object of the verb *bake* in our corpus 4 times.


- ▶ We can now compute *semantic similarity* in terms of *spatial distance*.
- ▶ One standard metric for this is the *Euclidean distance*:


$$d(\vec{a}, \vec{b}) = \sqrt{\sum_{i=1}^n (\vec{a}_i - \vec{b}_i)^2}$$

- ▶ Computes the norm (or *length*) of the *difference* of the vectors.
- ▶ The norm of a vector is:

$$\|\vec{x}\| = \sqrt{\sum_{i=1}^n x_i^2} = \sqrt{\vec{x} \cdot \vec{x}}$$


- ▶ Intuitive interpretation: The distance between two points corresponds to the length of the straight line connecting them.





- ▶ a: automobile
- ▶ b: car
- ▶ c: road
- ▶  $d(\vec{a}, \vec{b}) = 10$
- ▶  $d(\vec{a}, \vec{c}) = 7$

- ▶ However, a potential problem with Euclidean distance is that it is very sensitive to extreme values and the length of the vectors.
- ▶ As vectors of words with different *frequencies* will tend to have different length, the frequency will also affect the similarity judgment.




- ▶ One way to reduce frequency effects is to first **normalize** all our vectors to have **unit length**, i.e.  $\|\vec{x}\| = 1$
- ▶ Can be achieved by simply dividing each element by the length:  $\vec{x} \frac{1}{\|\vec{x}\|}$
- ▶ Amounts to all vectors pointing to the surface of a unit sphere.

- ▶ We can measure (cosine) *proximity* rather than (Euclidean) *distance*.
- ▶ Computes similarity as a function of the angle between the vectors:

$$\cos(\vec{a}, \vec{b}) = \frac{\sum_i \vec{a}_i \vec{b}_i}{\sqrt{\sum_i \vec{a}_i^2} \sqrt{\sum_i \vec{b}_i^2}} = \frac{\vec{a} \cdot \vec{b}}{\|\vec{a}\| \|\vec{b}\|}$$

- ▶ Constant range between 0 and 1.
- ▶ Avoids the arbitrary scaling caused by dimensionality, frequency, etc.
- ▶ As the angle between the vectors shortens, the cosine approaches 1.





- ▶ For *normalized* (unit) vectors, the cosine is simply the *dot product*:

$$\cos(\vec{a}, \vec{b}) = \vec{a} \cdot \vec{b} = \sum_{i=1}^n \vec{a}_i \vec{b}_i$$

- ▶ Can be computed very efficiently.
- ▶ The *same relative rank order* as the **Euclidean distance** for unit vectors!



- ▶ **Conceptually**, a vector space is often thought of as a **matrix**, often called **co-occurrence matrix** or word-context matrix.
  - ▶ Dimensions correspond to columns; each feature vector is a row.
  - ▶ For  $m$  words and  $n$  features we have an  $m \times n$  co-occurrence matrix.

## Corpus

- ▶ An **automobile** is a wheeled motor vehicle used for transporting passengers .
- ▶ A **car** is a form of transport, usually with four wheels and the capacity to carry around five passengers .
- ▶ Transport for the London games is limited , with spectators strongly advised to avoid the use of **cars** .

	advise	avoid	capacity	carry	...	vehicle	wheel	...
automobile	0	0	0	0	...	1	1	
car	1	1	1	1	...	0	1	



- ▶ As we move towards more realistic set-ups:
  - ▶ Semantic spaces will be extremely **high-dimensional**
  - ▶ The number of *non-zero* elements will be very low.
  - ▶ Few active features per word.
- ▶ We say that the vectors are **sparse**.
- ▶ This has implications for how to implement our data structures and vector operations:
- ▶ Don't want to waste space representing zero-valued features.



- ▶ In theory, you can view formulas like Euclidean norm and cosine as “pseudo-code” that you can translate directly into Lisp.
- ▶ But again; our feature vectors are sparse.
- ▶ Taken directly, a formula like the Euclidean norm requires iterating over every dimension  $n$  in our space.
- ▶ But we don't want to waste time iterating over zero elements if we don't have to!

- ▶ **Problem:** Raw co-occurrence frequencies are not very discriminative, and therefore not always the best indicators of relevance.
- ▶ Imagine we have some features recording information about direct objects and we've collected the following counts for the noun *wine*:
  - ▶  $\text{OBJ\_OF}(\text{buy}) = 14$
  - ▶  $\text{OBJ\_OF}(\text{pour}) = 8$
  - ▶ ... but the feature  $\text{OBJ\_OF}(\text{pour})$  seems more indicative of the semantics of *wine* than  $\text{OBJ\_OF}(\text{buy})$ .
- ▶ **Solution:** Weight the counts by an *association function*, “normalizing” our observed frequencies for chance co-occurrence.
- ▶ A range of different tests of statistical are used; e.g. *pointwise mutual information*, *log odds ratio*, *the t-test*, *log likelihood*, ...
- ▶ **Note:** We'll skip this step in our implementation (assignment 2a).

- ▶ So far we've looked at vector space models for detecting *words* with similar *meanings*.
- ▶ It's important to realize that vector space models are widely used for other purposes as well.
- ▶ Vector space models are commonly used in IR for finding *documents* with similar *content*.
- ▶ Each document  $d_j$  is represented by a feature vector, with features corresponding to the terms  $t_1, \dots, t_n$  occurring in the documents.
- ▶ Spatial distance  $\approx$  similarity of content.
- ▶ Can also represent a search *query* as a vector.
- ▶ The relevance of documents given by their distance to the query.

- ▶ The most commonly used weighting function is **tf-idf**:
  - ▶ **The term frequency**  $\text{tf}(t_i, d_j)$  denotes the number of times the term  $t_i$  occurs in document  $d_j$ .
  - ▶ **The document frequency**  $\text{df}(t_i)$  denotes the total number of documents in the collection that the term occurs in.
  - ▶ **The inverse document frequency** is defined as  $\text{idf}(t_i) = \log \left( \frac{N}{\text{df}(t_i)} \right)$ , where  $N$  is the total number of documents in the collection.
  - ▶ The weight given to term  $t_i$  in document  $d_j$  is then computed as

$$\text{tf-idf}(t_i, d_j) = \text{tf}(t_i, d_j) \times \text{idf}(t_i)$$

- ▶ A high tf-idf is obtained if a term has a *high* frequency in the given *document* and a *low* frequency in the document *collection* as whole.
- ▶ The weights hence tend to filter out common terms.

- ▶ Word meaning can be represented as a vector characterized by  $n$  dimensions.
- ▶ The  $n$  dimensions of our feature vectors represent the contextual features we observe.
- ▶ Raw co-occurrence counts are good but not the best way to quantify relevance.
- ▶ Semantic similarity can be computed based on spatial distance and proximity.
- ▶ We need to be careful when deciding on a data structure to represent the co-occurrence matrix and when we implement vector operations.



- ▶ Computing neighbor relations in the semantic space
- ▶ Representing classes
- ▶ Representing class membership
- ▶ Classification algorithms: KNN-classification /  $c$ -means, etc.



- Firth, J. R. (1957). A synopsis of linguistic theory 1930–1955. In *Studies in linguistic analysis*. Philological Society, Oxford.
- Harris, Z. S. (1968). *Mathematical structures of language*. New York: Wiley.
- Wittgenstein, L. (1953). *Philosophical investigations*. Oxford: Blackwell.