

*INF4820: Algorithms for  
Artificial Intelligence and  
Natural Language Processing*

Hidden Markov Models

Murhaf Fares & Stephan Oepen

Language Technology Group (LTG)

October 27, 2016





- ▶ Basic probability theory: axioms, joint vs. conditional probability, independence, Bayes' Theorem;



- ▶ Basic probability theory: axioms, joint vs. conditional probability, independence, Bayes' Theorem;
- ▶ Previous context can help predict the next element of a sequence, for example words in a sentence;



- ▶ Basic probability theory: axioms, joint vs. conditional probability, independence, Bayes' Theorem;
- ▶ Previous context can help predict the next element of a sequence, for example words in a sentence;
- ▶ Rather than use the whole previous context, the **Markov** assumption says that the whole history can be approximated by the last  $n - 1$  elements;



- ▶ Basic probability theory: axioms, joint vs. conditional probability, independence, Bayes' Theorem;
- ▶ Previous context can help predict the next element of a sequence, for example words in a sentence;
- ▶ Rather than use the whole previous context, the **Markov** assumption says that the whole history can be approximated by the last  $n - 1$  elements;
- ▶ An  $n$ -gram language model predicts the  $n$ -th word, conditioned on the  $n - 1$  previous words;





- ▶ Basic probability theory: axioms, joint vs. conditional probability, independence, Bayes' Theorem;
- ▶ Previous context can help predict the next element of a sequence, for example words in a sentence;
- ▶ Rather than use the whole previous context, the **Markov** assumption says that the whole history can be approximated by the last  $n - 1$  elements;
- ▶ An  $n$ -gram language model predicts the  $n$ -th word, conditioned on the  $n - 1$  previous words;
- ▶ Maximum Likelihood Estimation uses relative frequencies to approximate the conditional probabilities needed for an  $n$ -gram model;



- ▶ Basic probability theory: axioms, joint vs. conditional probability, independence, Bayes' Theorem;
- ▶ Previous context can help predict the next element of a sequence, for example words in a sentence;
- ▶ Rather than use the whole previous context, the **Markov** assumption says that the whole history can be approximated by the last  $n - 1$  elements;
- ▶ An  $n$ -gram language model predicts the  $n$ -th word, conditioned on the  $n - 1$  previous words;
- ▶ Maximum Likelihood Estimation uses relative frequencies to approximate the conditional probabilities needed for an  $n$ -gram model;
- ▶ Smoothing techniques are used to avoid zero probabilities.

## Determining


- ▶ which string is most likely:
  - ▶ *She studies morphosyntax* vs. *She studies more faux syntax*
- ▶ which tag sequence is most likely for *flies like flowers*:
  - ▶ **NNS VB NNS** vs. **VBZ P NNS**
- ▶ which syntactic analysis is most likely:






## Determining

- ▶ which string is most likely: ✓
  - ▶ *She studies morphosyntax* vs. *She studies more faux syntax*
- ▶ which tag sequence is most likely for *flies like flowers*:
  - ▶ **NNS VB NNS** vs. **VBZ P NNS**
- ▶ which syntactic analysis is most likely:



## Determining

- ▶ which string is most likely: ✓
  - ▶ *She studies morphosyntax* vs. *She studies more faux syntax*
- ▶ which tag sequence is most likely for *flies like flowers*:
  - ▶ **NNS VB NNS** vs. **VBZ P NNS**
- ▶ which syntactic analysis is most likely:



- ▶ Known by a variety of names: part-of-speech, POS, lexical categories, word classes, morphological classes, . . .
- ▶ ‘Traditionally’ defined semantically (e.g. “nouns are naming words”), but more accurately by their distributional properties.

- ▶ Known by a variety of names: part-of-speech, POS, lexical categories, word classes, morphological classes, . . .
- ▶ ‘Traditionally’ defined semantically (e.g. “nouns are naming words”), but more accurately by their distributional properties.
- ▶ Open-classes
  - ▶ New words created/updated/deleted all the time

- ▶ Known by a variety of names: part-of-speech, POS, lexical categories, word classes, morphological classes, . . .
- ▶ ‘Traditionally’ defined semantically (e.g. “nouns are naming words”), but more accurately by their distributional properties.
- ▶ Open-classes
  - ▶ New words created/updated/deleted all the time
- ▶ Closed-classes
  - ▶ Smaller classes, relatively static membership
  - ▶ Usually function words

- ▶ Nouns: dog, Oslo, scissors, snow, people, truth, cups
  - ▶ proper or common; countable or uncountable; plural or singular; masculine, feminine or neuter; . . .

- ▶ Nouns: dog, Oslo, scissors, snow, people, truth, cups
  - ▶ proper or common; countable or uncountable; plural or singular; masculine, feminine or neuter; . . .
- ▶ Verbs: fly, rained, having, ate, seen
  - ▶ transitive, intransitive, ditransitive; past, present, passive; stative or dynamic; plural or singular; . . .

- ▶ Nouns: dog, Oslo, scissors, snow, people, truth, cups
  - ▶ proper or common; countable or uncountable; plural or singular; masculine, feminine or neuter; . . .
- ▶ Verbs: fly, rained, having, ate, seen
  - ▶ transitive, intransitive, ditransitive; past, present, passive; stative or dynamic; plural or singular; . . .
- ▶ Adjectives: good, smaller, unique, fastest, best, unhappy
  - ▶ comparative or superlative; predicative or attributive; intersective or non-intersective; . . .



- ▶ Nouns: dog, Oslo, scissors, snow, people, truth, cups
  - ▶ proper or common; countable or uncountable; plural or singular; masculine, feminine or neuter; . . .
- ▶ Verbs: fly, rained, having, ate, seen
  - ▶ transitive, intransitive, ditransitive; past, present, passive; stative or dynamic; plural or singular; . . .
- ▶ Adjectives: good, smaller, unique, fastest, best, unhappy
  - ▶ comparative or superlative; predicative or attributive; intersective or non-intersective; . . .
- ▶ Adverbs: again, somewhat, slowly, yesterday, aloud
  - ▶ intersective; scopal; discourse; degree; temporal; directional; comparative or superlative; . . .

- ▶ Prepositions: *on, under, from, at, near, over, ...*
- ▶ Determiners: *a, an, the, that, ...*
- ▶ Pronouns: *she, who, I, others, ...*
- ▶ Conjunctions: *and, but, or, when, ...*
- ▶ Auxiliary verbs: *can, may, should, must, ...*
- ▶ Interjections, particles, numerals, negatives, politeness markers, greetings, existential there ...

(Examples from Jurafsky & Martin, 2008)

The (automatic) assignment of POS tags to word sequences

The (automatic) assignment of POS tags to word sequences

- ▶ non-trivial where words are ambiguous: *fly* (v) vs. *fly* (n)
- ▶ choice of the correct tag is *context-dependent*

The (automatic) assignment of POS tags to word sequences

- ▶ non-trivial where words are ambiguous: *fly* (v) vs. *fly* (n)
- ▶ choice of the correct tag is *context-dependent*
- ▶ useful in pre-processing for parsing, etc; but also directly for text-to-speech (TTS) system: **content** (n) vs. **content** (adj)

The (automatic) assignment of POS tags to word sequences

- ▶ non-trivial where words are ambiguous: *fly* (v) vs. *fly* (n)
  - ▶ choice of the correct tag is *context-dependent*
  - ▶ useful in pre-processing for parsing, etc; but also directly for text-to-speech (TTS) system: **content** (n) vs. **content** (adj)
  - ▶ difficulty and usefulness can depend on the *tagset*
    - ▶ English
      - ▶ Penn Treebank (PTB)—45 tags: NNS, NN, NNP, JJ, JJR, JJS
- <http://bulba.sdsu.edu/jeanette/thesis/PennTags.html>

The (automatic) assignment of POS tags to word sequences

- ▶ non-trivial where words are ambiguous: *fly* (v) vs. *fly* (n)
- ▶ choice of the correct tag is *context-dependent*
- ▶ useful in pre-processing for parsing, etc; but also directly for text-to-speech (TTS) system: **content** (n) vs. **content** (adj)
- ▶ difficulty and usefulness can depend on the *tagset*
  - ▶ English
    - ▶ Penn Treebank (PTB)—45 tags: NNS, NN, NNP, JJ, JJR, JJS  
<http://bulba.sdsu.edu/jeanette/thesis/PennTags.html>
  - ▶ Norwegian
    - ▶ Oslo-Bergen Tagset—multi-part: ⟨subst appell fem be ent⟩  
<http://tekstlab.uio.no/obt-ny/english/tags.html>

- We are interested in the probability of sequences like:

flies	like	the	wind	or	flies	like	the	wind
NNS	VB	DT	NN		VBZ	P	DT	NN



- We are interested in the probability of sequences like:

flies	like	the	wind	or	flies	like	the	wind
NNS	VB	DT	NN		VBZ	P	DT	NN

- In normal text, we see the words, but not the tags.

- We are interested in the probability of sequences like:

flies	like	the	wind	or	flies	like	the	wind
NNS	VB	DT	NN		VBZ	P	DT	NN

- In normal text, we see the words, but not the tags.
- Consider the POS tags to be underlying skeleton of the sentence, unseen but influencing the sentence shape.

- ▶ We are interested in the probability of sequences like:

flies	like	the	wind	or	flies	like	the	wind
NNS	VB	DT	NN		VBZ	P	DT	NN


- ▶ In normal text, we see the words, but not the tags.
- ▶ Consider the POS tags to be underlying skeleton of the sentence, unseen but influencing the sentence shape.
- ▶ A structure like this, consisting of a **hidden** state sequence, and a related **observation** sequence can be modelled as a *Hidden Markov Model*.

The generative story:




$\langle S \rangle$

The generative story:




$$P(S, O) = P(\text{DT}|\langle S \rangle)$$

The generative story:




$$P(S, O) = P(\text{DT}|\langle S \rangle) P(\text{the}|\text{DT})$$

The generative story:



$$P(S, O) = P(\text{DT}|\langle S \rangle) P(\text{the}|\text{DT}) P(\text{NN}|\text{DT})$$


The generative story:



$$P(S, O) = P(DT|\langle S \rangle) P(\text{the}|DT) P(NN|DT) P(\text{cat}|NN)$$




The generative story:




$$P(S, O) = P(\text{DT}|\langle S \rangle) P(\text{the}|\text{DT}) P(\text{NN}|\text{DT}) P(\text{cat}|\text{NN}) \\ P(\text{VBZ}|\text{NN})$$

The generative story:




$$P(S, O) = P(\text{DT}|\langle S \rangle) P(\text{the}|\text{DT}) P(\text{NN}|\text{DT}) P(\text{cat}|\text{NN}) \\ P(\text{VBZ}|\text{NN}) P(\text{eats}|\text{VBZ})$$

The generative story:




$$P(S, O) = P(\text{DT}|\langle S \rangle) P(\text{the}|\text{DT}) P(\text{NN}|\text{DT}) P(\text{cat}|\text{NN}) \\ P(\text{VBZ}|\text{NN}) P(\text{eats}|\text{VBZ}) P(\text{NNS}|\text{VBZ})$$

The generative story:




$$P(S, O) = P(\text{DT}|\langle S \rangle) P(\text{the}|\text{DT}) P(\text{NN}|\text{DT}) P(\text{cat}|\text{NN}) \\ P(\text{VBZ}|\text{NN}) P(\text{eats}|\text{VBZ}) P(\text{NNS}|\text{VBZ}) P(\text{mice}|\text{NNS})$$

The generative story:



$$\begin{aligned} P(S, O) = & P(\text{DT} | \langle S \rangle) P(\text{the} | \text{DT}) P(\text{NN} | \text{DT}) P(\text{cat} | \text{NN}) \\ & P(\text{VBZ} | \text{NN}) P(\text{eats} | \text{VBZ}) P(\text{NNS} | \text{VBZ}) P(\text{mice} | \text{NNS}) \\ & P(\langle /S \rangle | \text{NNS}) \end{aligned}$$

The generative story:



$$\begin{aligned} P(S, O) = & P(\text{DT}|\langle S \rangle) P(\text{the}|\text{DT}) P(\text{NN}|\text{DT}) P(\text{cat}|\text{NN}) \\ & P(\text{VBZ}|\text{NN}) P(\text{eats}|\text{VBZ}) P(\text{NNS}|\text{VBZ}) P(\text{mice}|\text{NNS}) \\ & P(\langle /S \rangle|\text{NNS}) \end{aligned}$$

For a bi-gram HMM, with  $O_1^N$ :

$$P(S, O) = \prod_{i=1}^{N+1} P(s_i | s_{i-1}) P(o_i | s_i) \quad \text{where} \quad s_0 = \langle S \rangle, s_{N+1} = \langle /S \rangle$$

For a bi-gram HMM, with  $O_1^N$ :

$$P(S, O) = \prod_{i=1}^{N+1} P(s_i | s_{i-1}) P(o_i | s_i) \quad \text{where} \quad s_0 = \langle S \rangle, s_{N+1} = \langle /S \rangle$$

- The **transition probabilities** model the probabilities of moving from state to state.



For a bi-gram HMM, with  $O_1^N$ :

$$P(S, O) = \prod_{i=1}^{N+1} P(s_i | s_{i-1}) P(o_i | s_i) \quad \text{where} \quad s_0 = \langle S \rangle, s_{N+1} = \langle /S \rangle$$

- ▶ The transition probabilities model the probabilities of moving from state to state.
- ▶ The **emission probabilities** model the probability that a state *emits* a particular observation.

The HMM models the process of generating the labelled sequence. We can use this model for a number of tasks:

- ▶  $P(S, O)$  given  $S$  and  $O$
- ▶  $P(O)$  given  $O$
- ▶  $S$  that maximises  $P(S|O)$  given  $O$
- ▶ We can also learn the model parameters, given a set of observations.

The HMM models the process of generating the labelled sequence. We can use this model for a number of tasks:

- ▶  $P(S, O)$  given  $S$  and  $O$
- ▶  $P(O)$  given  $O$
- ▶  $S$  that maximises  $P(S|O)$  given  $O$
- ▶ We can also learn the model parameters, given a set of observations.

The HMM models the process of generating the labelled sequence. We can use this model for a number of tasks:

- ▶  $P(S, O)$  given  $S$  and  $O$
- ▶  $P(O)$  given  $O$
- ▶  $S$  that maximises  $P(S|O)$  given  $O$
- ▶ We can also learn the model parameters, given a set of observations.



The HMM models the process of generating the labelled sequence. We can use this model for a number of tasks:

- ▶  $P(S, O)$  given  $S$  and  $O$
- ▶  $P(O)$  given  $O$
- ▶  $S$  that maximises  $P(S|O)$  given  $O$
- ▶ We can also learn the model parameters, given a set of observations.

Our *observations* will be **words** ( $w_i$ ), and our *states* PoS **tags** ( $t_i$ )

As so often in NLP, we learn an HMM from labelled data:

## Transition probabilities

Based on a training corpus of previously tagged text, with tags as our state, the MLE can be computed from the counts of observed tags:

$$P(t_i|t_{i-1}) = \frac{C(t_{i-1}, t_i)}{C(t_{i-1})}$$

As so often in NLP, we learn an HMM from labelled data:

## Transition probabilities

Based on a training corpus of previously tagged text, with tags as our state, the MLE can be computed from the counts of observed tags:

$$P(t_i|t_{i-1}) = \frac{C(t_{i-1}, t_i)}{C(t_{i-1})}$$

## Emission probabilities

Computed from relative frequencies in the same way, with the words as observations:

$$P(w_i|t_i) = \frac{C(t_i, w_i)}{C(t_i)}$$



$$\begin{aligned}P(S, O) &= P(s_1|\langle S \rangle)P(o_1|s_1)P(s_2|s_1)P(o_2|s_2)P(s_3|s_2)P(o_3|s_3) \dots \\&= 0.0429 \times 0.0031 \times 0.0044 \times 0.0001 \times 0.0072 \times \dots\end{aligned}$$



$$\begin{aligned}P(S, O) &= P(s_1|\langle S \rangle)P(o_1|s_1)P(s_2|s_1)P(o_2|s_2)P(s_3|s_2)P(o_3|s_3) \dots \\&= 0.0429 \times 0.0031 \times 0.0044 \times 0.0001 \times 0.0072 \times \dots\end{aligned}$$

- Multiplying many small probabilities  $\rightarrow$  underflow

$$\begin{aligned}P(S, O) &= P(s_1|\langle S \rangle)P(o_1|s_1)P(s_2|s_1)P(o_2|s_2)P(s_3|s_2)P(o_3|s_3) \dots \\&= 0.0429 \times 0.0031 \times 0.0044 \times 0.0001 \times 0.0072 \times \dots\end{aligned}$$

- ▶ Multiplying many small probabilities  $\rightarrow$  underflow
- ▶ Solution: work in log(arithmetic) space:
  - ▶  $\log(AB) = \log(A) + \log(B)$
  - ▶ hence  $P(A)P(B) = \exp(\log(A) + \log(B))$
  - ▶  $\log(P(S, O)) = -1.368 + -2.509 + -2.357 + -4 + -2.143 + \dots$

$$\begin{aligned}P(S, O) &= P(s_1|\langle S \rangle)P(o_1|s_1)P(s_2|s_1)P(o_2|s_2)P(s_3|s_2)P(o_3|s_3) \dots \\&= 0.0429 \times 0.0031 \times 0.0044 \times 0.0001 \times 0.0072 \times \dots\end{aligned}$$

- ▶ Multiplying many small probabilities  $\rightarrow$  underflow
- ▶ Solution: work in log(arithmetic) space:
  - ▶  $\log(AB) = \log(A) + \log(B)$
  - ▶ hence  $P(A)P(B) = \exp(\log(A) + \log(B))$
  - ▶  $\log(P(S, O)) = -1.368 + -2.509 + -2.357 + -4 + -2.143 + \dots$

The issues related to MLE / smoothing that we discussed for  $n$ -gram models also applies here ...

## Missing records of weather in Baltimore for Summer 2007

- ▶ Jason likes to eat ice cream.
- ▶ He records his daily ice cream consumption in his diary.
- ▶ The number of ice creams he ate was influenced, but not entirely determined by the weather.
- ▶ Today's weather is partially predictable from yesterday's.

## Missing records of weather in Baltimore for Summer 2007


- ▶ Jason likes to eat ice cream.
- ▶ He records his daily ice cream consumption in his diary.
- ▶ The number of ice creams he ate was influenced, but not entirely determined by the weather.
- ▶ Today's weather is partially predictable from yesterday's.

## A Hidden Markov Model!

with:

- ▶ Hidden states:  $\{H, C\}$  (plus pseudo-states  $\langle S \rangle$  and  $\langle /S \rangle$ )
- ▶ Observations:  $\{1, 2, 3\}$

# Ice Cream and Global Warming



The HMM models the process of generating the labelled sequence. We can use this model for a number of tasks:

- ▶  $P(S, O)$  given  $S$  and  $O$
- ▶  $P(O)$  given  $O$
- ▶  $S$  that maximises  $P(S|O)$  given  $O$
- ▶  $P(s_x|O)$  given  $O$
- ▶ We can also learn the model parameters, given a set of observations.

We want to find the tag sequence, given a word sequence. With tags as our states and words as our observations, we know:

$$P(S, O) = \prod_{i=1}^{N+1} P(s_i | s_{i-1}) P(o_i | s_i)$$

We want:  $P(S|O)$



We want to find the tag sequence, given a word sequence. With tags as our states and words as our observations, we know:

$$P(S, O) = \prod_{i=1}^{N+1} P(s_i | s_{i-1}) P(o_i | s_i)$$

We want:  $P(S|O) = \frac{P(S, O)}{P(O)}$

We want to find the tag sequence, given a word sequence. With tags as our states and words as our observations, we know:

$$P(S, O) = \prod_{i=1}^{N+1} P(s_i | s_{i-1}) P(o_i | s_i)$$

We want:  $P(S|O) = \frac{P(S, O)}{P(O)}$

Actually, we want the state sequence  $\hat{S}$  that maximises  $P(S|O)$ :

$$\hat{S} = \arg \max_S \frac{P(S, O)}{P(O)}$$

We want to find the tag sequence, given a word sequence. With tags as our states and words as our observations, we know:

$$P(S, O) = \prod_{i=1}^{N+1} P(s_i | s_{i-1}) P(o_i | s_i)$$

We want:  $P(S|O) = \frac{P(S, O)}{P(O)}$

Actually, we want the state sequence  $\hat{S}$  that maximises  $P(S|O)$ :

$$\hat{S} = \arg \max_S \frac{P(S, O)}{P(O)}$$

Since  $P(O)$  always is the same, we can drop the denominator:

$$\hat{S} = \arg \max_S P(S, O)$$

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM

$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$
$P(H H) = 0.6$	$P(C H) = 0.2$
$P(H C) = 0.3$	$P(C C) = 0.5$
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$
$P(1 H) = 0.2$	$P(1 C) = 0.5$
$P(2 H) = 0.4$	$P(2 C) = 0.4$
$P(3 H) = 0.4$	$P(3 C) = 0.1$

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM

if  $O = 3\ 1\ 3$

$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$
$P(H H) = 0.6$	$P(C H) = 0.2$
$P(H C) = 0.3$	$P(C C) = 0.5$
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$
$P(1 H) = 0.2$	$P(1 C) = 0.5$
$P(2 H) = 0.4$	$P(2 C) = 0.4$
$P(3 H) = 0.4$	$P(3 C) = 0.1$

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM

if  $O = 3\ 1\ 3$

$$P(H|\langle S \rangle) = 0.8 \quad P(C|\langle S \rangle) = 0.2$$

$$P(H|H) = 0.6 \quad P(C|H) = 0.2$$

$$P(H|C) = 0.3 \quad P(C|C) = 0.5$$

$$P(\langle /S \rangle | H) = 0.2 \quad P(\langle /S \rangle | C) = 0.2$$

$$P(1|H) = 0.2 \quad P(1|C) = 0.5$$

$$P(2|H) = 0.4 \quad P(2|C) = 0.4$$

$$P(3|H) = 0.4 \quad P(3|C) = 0.1$$

$\langle S \rangle$     H    H    H     $\langle /S \rangle$

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM

if  $O = 3\ 1\ 3$

$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$
$P(H H) = 0.6$	$P(C H) = 0.2$
$P(H C) = 0.3$	$P(C C) = 0.5$
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$
$P(1 H) = 0.2$	$P(1 C) = 0.5$
$P(2 H) = 0.4$	$P(2 C) = 0.4$
$P(3 H) = 0.4$	$P(3 C) = 0.1$

$\langle S \rangle$	H	H	H	$\langle /S \rangle$	0.0018432
---------------------	---	---	---	----------------------	-----------

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM

if  $O = 3\ 1\ 3$

$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$
$P(H H) = 0.6$	$P(C H) = 0.2$
$P(H C) = 0.3$	$P(C C) = 0.5$
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$
$P(1 H) = 0.2$	$P(1 C) = 0.5$
$P(2 H) = 0.4$	$P(2 C) = 0.4$
$P(3 H) = 0.4$	$P(3 C) = 0.1$

$\langle S \rangle$	H	H	H	$\langle /S \rangle$	0.0018432
$\langle S \rangle$	H	H	C	$\langle /S \rangle$	



## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM

if  $O = 3 \ 1 \ 3$

$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$
$P(H H) = 0.6$	$P(C H) = 0.2$
$P(H C) = 0.3$	$P(C C) = 0.5$
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$
$P(1 H) = 0.2$	$P(1 C) = 0.5$
$P(2 H) = 0.4$	$P(2 C) = 0.4$
$P(3 H) = 0.4$	$P(3 C) = 0.1$

$\langle S \rangle$	H	H	H	$\langle /S \rangle$	0.0018432
$\langle S \rangle$	H	H	C	$\langle /S \rangle$	0.0001536

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM

if  $O = 3\ 1\ 3$

$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$
$P(H H) = 0.6$	$P(C H) = 0.2$
$P(H C) = 0.3$	$P(C C) = 0.5$
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$
$P(1 H) = 0.2$	$P(1 C) = 0.5$
$P(2 H) = 0.4$	$P(2 C) = 0.4$
$P(3 H) = 0.4$	$P(3 C) = 0.1$

$\langle S \rangle$	H	H	H	$\langle /S \rangle$	0.0018432
$\langle S \rangle$	H	H	C	$\langle /S \rangle$	0.0001536
$\langle S \rangle$	H	C	H	$\langle /S \rangle$	0.0007680

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM

if  $O = 3\ 1\ 3$

$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$
$P(H H) = 0.6$	$P(C H) = 0.2$
$P(H C) = 0.3$	$P(C C) = 0.5$
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$
$P(1 H) = 0.2$	$P(1 C) = 0.5$
$P(2 H) = 0.4$	$P(2 C) = 0.4$
$P(3 H) = 0.4$	$P(3 C) = 0.1$

$\langle S \rangle$	H	H	H	$\langle /S \rangle$	0.0018432
$\langle S \rangle$	H	H	C	$\langle /S \rangle$	0.0001536
$\langle S \rangle$	H	C	H	$\langle /S \rangle$	0.0007680
$\langle S \rangle$	H	C	C	$\langle /S \rangle$	0.0003200

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM		if $O = 3\ 1\ 3$				
$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$	$\langle S \rangle$	H	H	H	$\langle /S \rangle$ 0.0018432
$P(H H) = 0.6$	$P(C H) = 0.2$	$\langle S \rangle$	H	H	C	$\langle /S \rangle$ 0.0001536
$P(H C) = 0.3$	$P(C C) = 0.5$	$\langle S \rangle$	H	C	H	$\langle /S \rangle$ 0.0007680
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$	$\langle S \rangle$	H	C	C	$\langle /S \rangle$ 0.0003200
$P(1 H) = 0.2$	$P(1 C) = 0.5$	$\langle S \rangle$	C	H	H	$\langle /S \rangle$ 0.0000576
$P(2 H) = 0.4$	$P(2 C) = 0.4$					
$P(3 H) = 0.4$	$P(3 C) = 0.1$					

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM		if $O = 3\ 1\ 3$				
$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$	$\langle S \rangle$	H	H	H	$\langle /S \rangle$ 0.0018432
$P(H H) = 0.6$	$P(C H) = 0.2$	$\langle S \rangle$	H	H	C	$\langle /S \rangle$ 0.0001536
$P(H C) = 0.3$	$P(C C) = 0.5$	$\langle S \rangle$	H	C	H	$\langle /S \rangle$ 0.0007680
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$	$\langle S \rangle$	H	C	C	$\langle /S \rangle$ 0.0003200
$P(1 H) = 0.2$	$P(1 C) = 0.5$	$\langle S \rangle$	C	H	H	$\langle /S \rangle$ 0.0000576
$P(2 H) = 0.4$	$P(2 C) = 0.4$	$\langle S \rangle$	C	H	C	$\langle /S \rangle$ 0.0000048
$P(3 H) = 0.4$	$P(3 C) = 0.1$					

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM		if $O = 3\ 1\ 3$				
$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$	$\langle S \rangle$	H	H	H	$\langle /S \rangle$ 0.0018432
$P(H H) = 0.6$	$P(C H) = 0.2$	$\langle S \rangle$	H	H	C	$\langle /S \rangle$ 0.0001536
$P(H C) = 0.3$	$P(C C) = 0.5$	$\langle S \rangle$	H	C	H	$\langle /S \rangle$ 0.0007680
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$	$\langle S \rangle$	H	C	C	$\langle /S \rangle$ 0.0003200
$P(1 H) = 0.2$	$P(1 C) = 0.5$	$\langle S \rangle$	C	H	H	$\langle /S \rangle$ 0.0000576
$P(2 H) = 0.4$	$P(2 C) = 0.4$	$\langle S \rangle$	C	H	C	$\langle /S \rangle$ 0.0000048
$P(3 H) = 0.4$	$P(3 C) = 0.1$	$\langle S \rangle$	C	C	H	$\langle /S \rangle$ 0.0001200

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM		if $O = 3\ 1\ 3$					
$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$	$\langle S \rangle$	H	H	H	$\langle /S \rangle$	0.0018432
$P(H H) = 0.6$	$P(C H) = 0.2$	$\langle S \rangle$	H	H	C	$\langle /S \rangle$	0.0001536
$P(H C) = 0.3$	$P(C C) = 0.5$	$\langle S \rangle$	H	C	H	$\langle /S \rangle$	0.0007680
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$	$\langle S \rangle$	H	C	C	$\langle /S \rangle$	0.0003200
$P(1 H) = 0.2$	$P(1 C) = 0.5$	$\langle S \rangle$	C	H	H	$\langle /S \rangle$	0.0000576
$P(2 H) = 0.4$	$P(2 C) = 0.4$	$\langle S \rangle$	C	H	C	$\langle /S \rangle$	0.0000048
$P(3 H) = 0.4$	$P(3 C) = 0.1$	$\langle S \rangle$	C	C	H	$\langle /S \rangle$	0.0001200
		$\langle S \rangle$	C	C	C	$\langle /S \rangle$	0.0000500

## Task

What is the most likely state sequence  $S$ , given an observation sequence  $O$  and an HMM.

HMM		if $O = 3\ 1\ 3$					
$P(H \langle S \rangle) = 0.8$	$P(C \langle S \rangle) = 0.2$	$\langle S \rangle$	H	H	H	$\langle /S \rangle$	0.0018432
$P(H H) = 0.6$	$P(C H) = 0.2$	$\langle S \rangle$	H	H	C	$\langle /S \rangle$	0.0001536
$P(H C) = 0.3$	$P(C C) = 0.5$	$\langle S \rangle$	H	C	H	$\langle /S \rangle$	0.0007680
$P(\langle /S \rangle H) = 0.2$	$P(\langle /S \rangle C) = 0.2$	$\langle S \rangle$	H	C	C	$\langle /S \rangle$	0.0003200
$P(1 H) = 0.2$	$P(1 C) = 0.5$	$\langle S \rangle$	C	H	H	$\langle /S \rangle$	0.0000576
$P(2 H) = 0.4$	$P(2 C) = 0.4$	$\langle S \rangle$	C	H	C	$\langle /S \rangle$	0.0000048
$P(3 H) = 0.4$	$P(3 C) = 0.1$	$\langle S \rangle$	C	C	H	$\langle /S \rangle$	0.0001200
		$\langle S \rangle$	C	C	C	$\langle /S \rangle$	0.0000500



For (only) two states and a (short) observation sequence of length three, comparing all possible sequences is workable, but . . .

For (only) two states and a (short) observation sequence of length three, comparing all possible sequences is workable, but . . .

- ▶ for  $N$  observations and  $L$  states, there are  $L^N$  sequences
- ▶ we do the same partial calculations over and over again

For (only) two states and a (short) observation sequence of length three, comparing all possible sequences is workable, but . . .

- ▶ for  $N$  observations and  $L$  states, there are  $L^N$  sequences
- ▶ we do the same partial calculations over and over again

## Dynamic Programming:

- ▶ records sub-problem solutions for further re-use
- ▶ useful when a complex problem can be described recursively
- ▶ examples: Dijkstra's shortest path, minimum edit distance, longest common subsequence, Viterbi algorithm

For (only) two states and a (short) observation sequence of length three, comparing all possible sequences is workable, but . . .

- ▶ for  $N$  observations and  $L$  states, there are  $L^N$  sequences
- ▶ we do the same partial calculations over and over again

## Dynamic Programming:

- ▶ records sub-problem solutions for further re-use
- ▶ useful when a complex problem can be described recursively
- ▶ examples: Dijkstra's shortest path, minimum edit distance, longest common subsequence, **Viterbi algorithm**

Recall our problem:

$$\text{maximise } P(s_1 \dots s_n | o_1 \dots o_n) = P(s_1 | s_0) P(o_1 | s_1) P(s_2 | s_1) P(o_2 | s_2) \dots$$

Recall our problem:

$$\text{maximise } P(s_1 \dots s_n | o_1 \dots o_n) = P(s_1 | s_0) P(o_1 | s_1) P(s_2 | s_1) P(o_2 | s_2) \dots$$

Our recursive sub-problem:

$$v_i(x) = \max_{k=1}^L [v_{i-1}(k) \cdot P(x|k) \cdot P(o_i|x)]$$

The variable  $v_i(x)$  represents the maximum probability that the  $i$ -th state is  $x$ , given that we have seen  $O_1^i$ .

Recall our problem:

$$\text{maximise } P(s_1 \dots s_n | o_1 \dots o_n) = P(s_1 | s_0) P(o_1 | s_1) P(s_2 | s_1) P(o_2 | s_2) \dots$$

Our **recursive** sub-problem:

$$v_i(x) = \max_{k=1}^L [v_{i-1}(k) \cdot P(x|k) \cdot P(o_i|x)]$$

The variable  $v_i(x)$  represents the maximum probability that the  $i$ -th state is  $x$ , given that we have seen  $O_1^i$ .

Recall our problem:

$$\text{maximise } P(s_1 \dots s_n | o_1 \dots o_n) = P(s_1 | s_0) P(o_1 | s_1) P(s_2 | s_1) P(o_2 | s_2) \dots$$

Our recursive sub-problem:


$$v_i(x) = \max_{k=1}^L [v_{i-1}(k) \cdot P(x|k) \cdot P(o_i|x)]$$

The variable  $v_i(x)$  represents the maximum probability that the  $i$ -th state is  $x$ , given that we have seen  $O_1^i$ .


At each step, we record backpointers showing which previous state led to the maximum probability.




# An Example of the Viterbi Algorithm




# An Example of the Viterbi Algorithm




# An Example of the Viterbi Algorithm



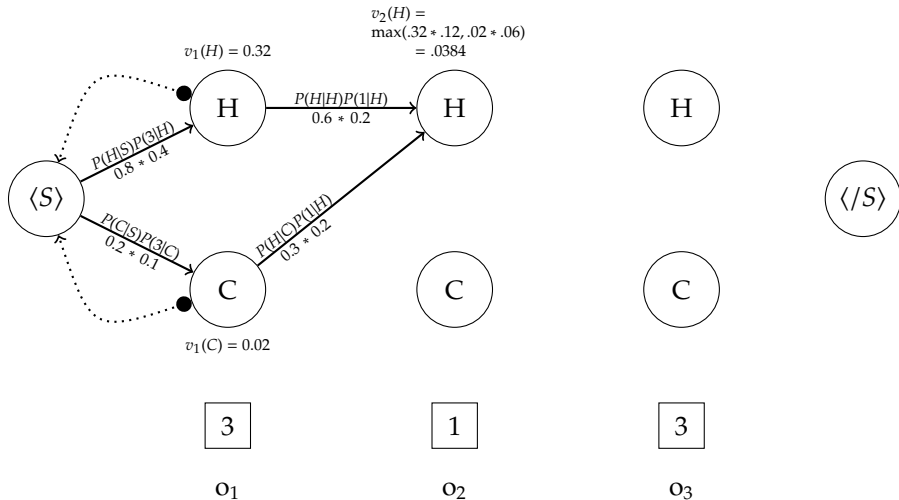
# An Example of the Viterbi Algorithm




# An Example of the Viterbi Algorithm



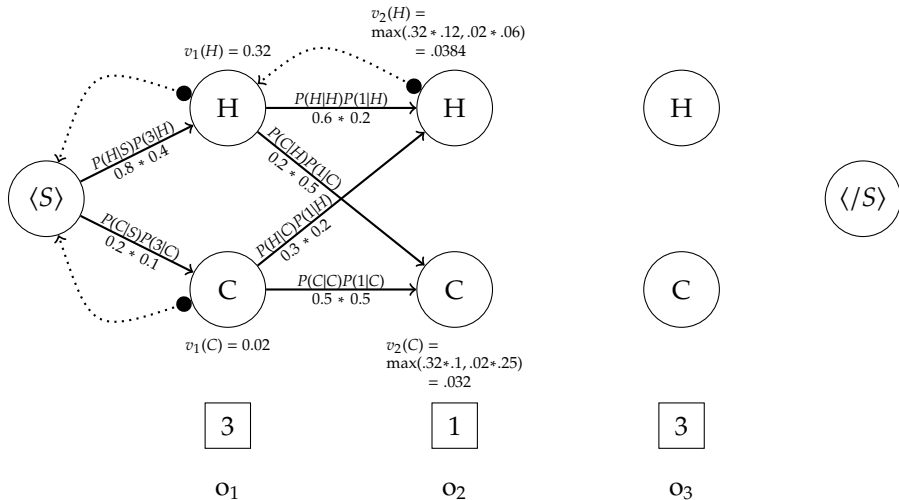
# An Example of the Viterbi Algorithm



# An Example of the Viterbi Algorithm

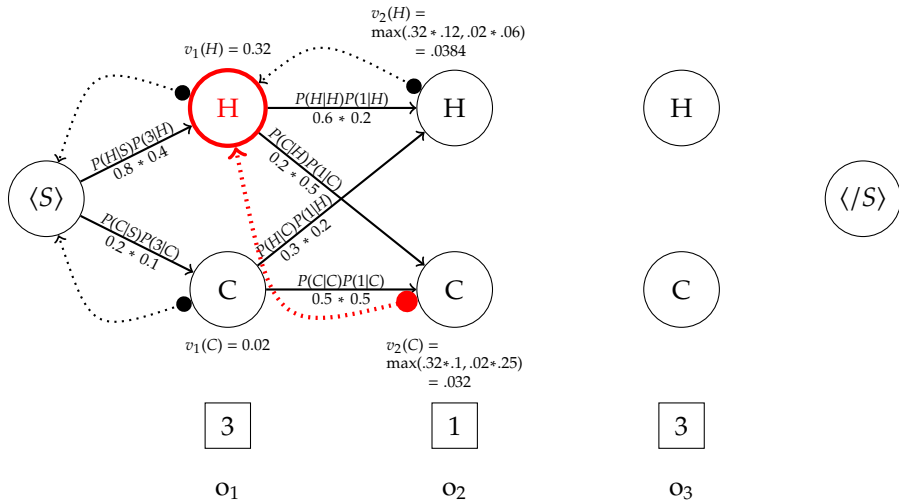


# An Example of the Viterbi Algorithm

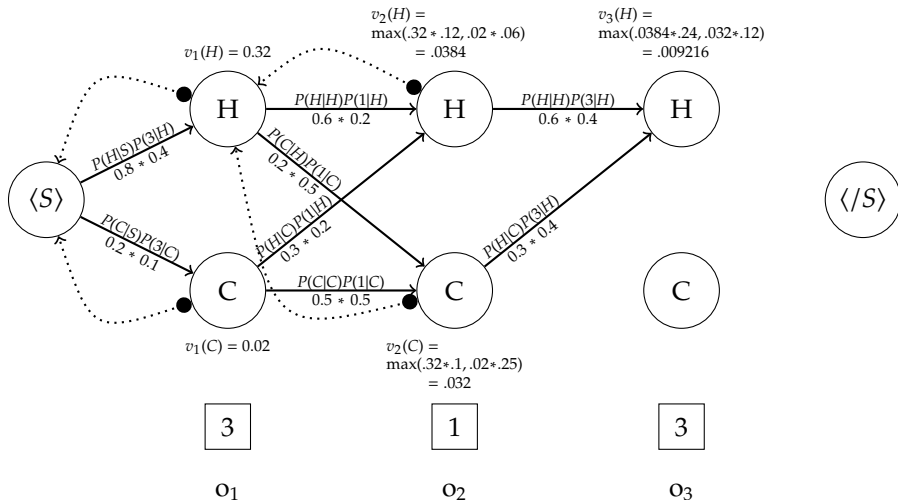




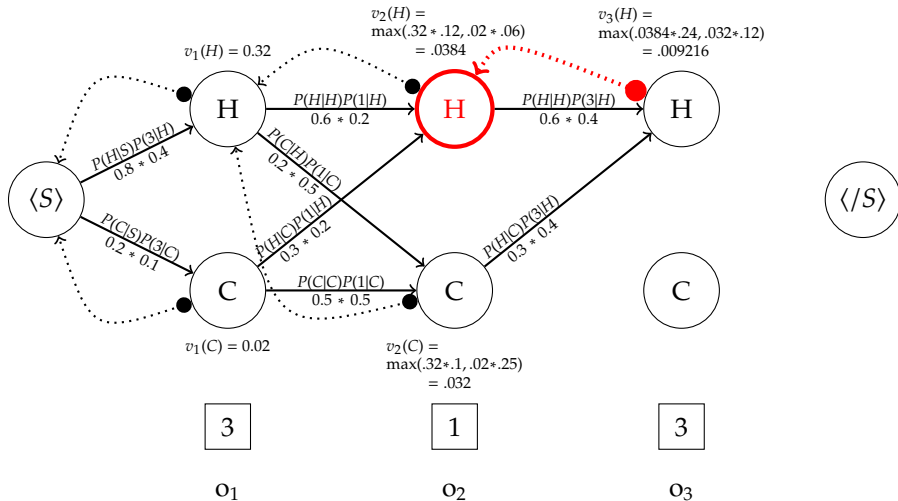
# An Example of the Viterbi Algorithm



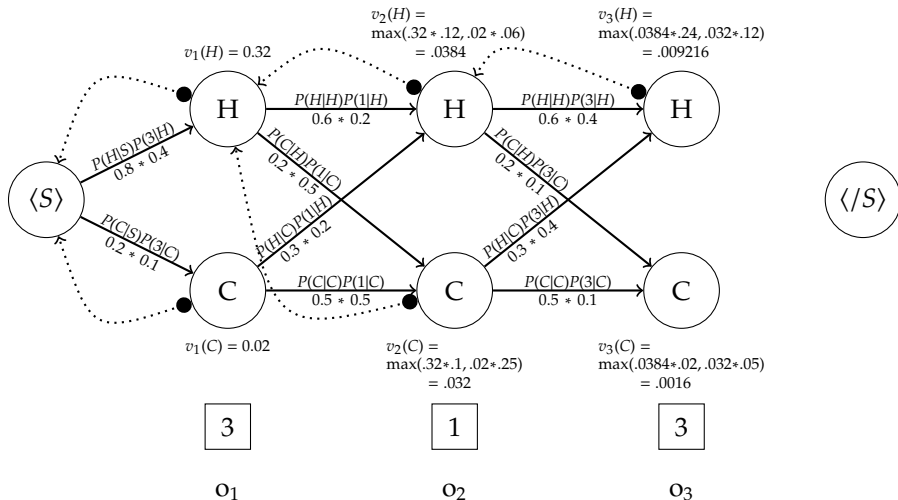
# An Example of the Viterbi Algorithm




# An Example of the Viterbi Algorithm



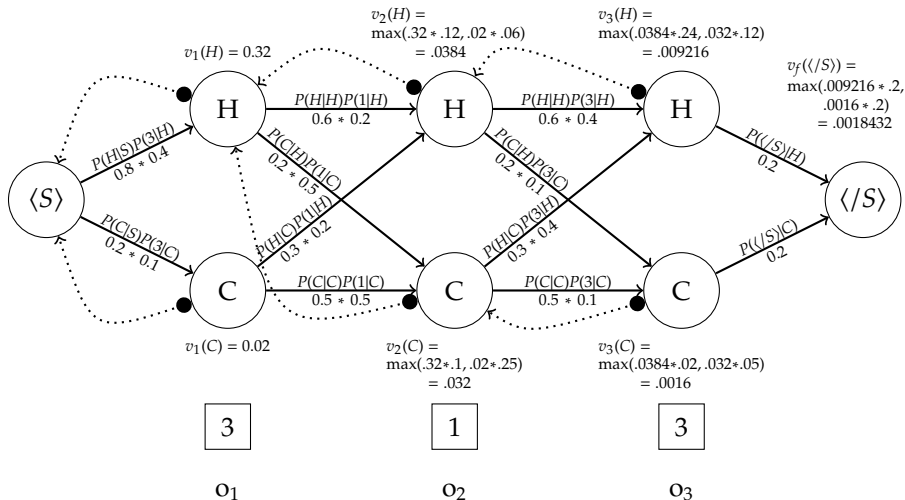
# An Example of the Viterbi Algorithm




# An Example of the Viterbi Algorithm



# An Example of the Viterbi Algorithm




# An Example of the Viterbi Algorithm




H

# An Example of the Viterbi Algorithm

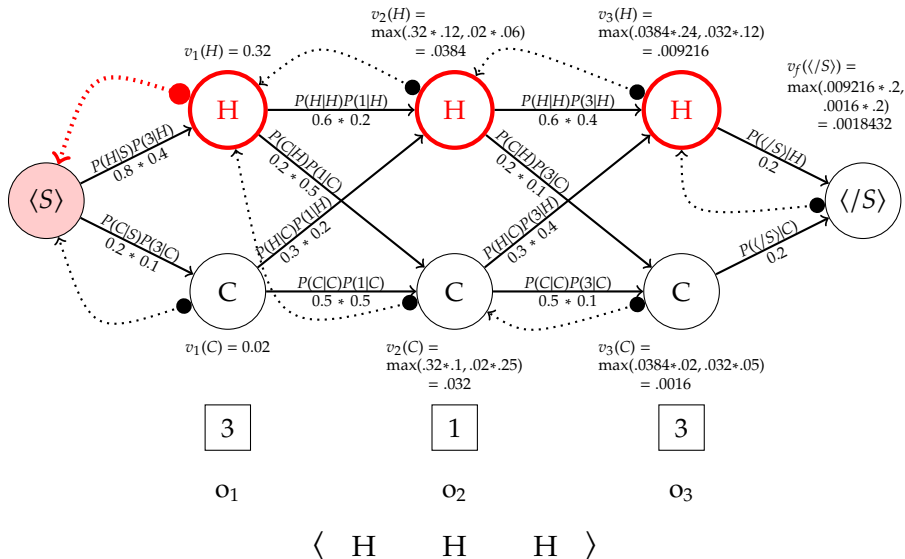




# An Example of the Viterbi Algorithm



# An Example of the Viterbi Algorithm



# Pseudocode for the Viterbi Algorithm



**Input:** *observations* of length  $N$ , state set of size  $L$

**Output:** *best-path*

create a path probability matrix  $viterbi[N, L + 2]$

create a path backpointer matrix  $backpointer[N, L + 2]$

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$

$backpointer[1, s] \leftarrow 0$

**end**

**for each** time step  $i$  from 2 to  $N$  **do**

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[i, s] \leftarrow \max_{s'=1}^L viterbi[i - 1, s'] \times trans(s', s) \times emit(o_i, s)$

$backpointer[i, s] \leftarrow \arg \max_{s'=1}^L viterbi[i - 1, s'] \times trans(s', s)$

**end**

**end**

$viterbi[N, L + 1] \leftarrow \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

$backpointer[N, L + 1] \leftarrow \arg \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

**return** the path by following backpointers from  $backpointer[N, L + 1]$



Big-O notation describes the complexity of an algorithm.

- ▶ it describes the worst-case *order of growth* in terms of the size of the input
- ▶ only the largest order term is represented
- ▶ constant factors are ignored
- ▶ determined by looking at loops in the code

# Pseudocode for the Viterbi Algorithm



**Input:** *observations* of length  $N$ , state set of length  $L$

**Output:** *best-path*

create a path probability matrix  $viterbi[N, L + 2]$

create a path backpointer matrix  $backpointer[N, L + 2]$

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$

$backpointer[1, s] \leftarrow 0$

**end**

**for each** time step  $i$  from 2 to  $N$  **do**

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[i, s] \leftarrow \max_{s'=1}^L viterbi[i - 1, s'] \times trans(s', s) \times emit(o_i, s)$

$backpointer[i, s] \leftarrow \arg \max_{s'=1}^L viterbi[i - 1, s'] \times trans(s', s)$

**end**

**end**

$viterbi[N, L + 1] \leftarrow \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

$backpointer[N, L + 1] \leftarrow \arg \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

**return** the path by following backpointers from  $backpointer[N, L + 1]$

# Pseudocode for the Viterbi Algorithm



**Input:** *observations* of length  $N$ , state set of length  $L$

**Output:** *best-path*

create a path probability matrix  $viterbi[N, L + 2]$

create a path backpointer matrix  $backpointer[N, L + 2]$

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$   
     $backpointer[1, s] \leftarrow 0$

**end**

**for each** time step  $i$  from 2 to  $N$  **do**

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[i, s] \leftarrow \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s) \times emit(o_i, s)$   
         $backpointer[i, s] \leftarrow \arg \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s)$

**end**

**end**

$viterbi[N, L + 1] \leftarrow \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

$backpointer[N, L + 1] \leftarrow \arg \max_{s=1}^L viterbi[N, s] \times trans(s, \langle /S \rangle)$

**return** the path by following backpointers from  $backpointer[N, L + 1]$

$L$

$L$

# Pseudocode for the Viterbi Algorithm



**Input:** observations of length  $N$ , state set of length  $L$

**Output:** best-path

create a path probability matrix  $viterbi[N, L + 2]$

create a path backpointer matrix  $backpointer[N, L + 2]$

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$

$backpointer[1, s] \leftarrow 0$

**end**

**for each** time step  $i$  from 2 to  $N$  **do**

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[i, s] \leftarrow \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s) \times emit(o_i, s)$

$backpointer[i, s] \leftarrow \arg \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s)$

**end**

**end**

$viterbi[N, L + 1] \leftarrow \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

$backpointer[N, L + 1] \leftarrow \arg \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

**return** the path by following backpointers from  $backpointer[N, L + 1]$

$L$

# Pseudocode for the Viterbi Algorithm



**Input:** observations of length  $N$ , state set of length  $L$

**Output:** best-path

create a path probability matrix  $viterbi[N, L + 2]$

create a path backpointer matrix  $backpointer[N, L + 2]$

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$

$backpointer[1, s] \leftarrow 0$

**end**

**for each** time step  $i$  from 2 to  $N$  **do**

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[i, s] \leftarrow \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s) \times emit(o_i, s)$

$backpointer[i, s] \leftarrow \arg \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s)$

**end**

**end**

$viterbi[N, L + 1] \leftarrow \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

$backpointer[N, L + 1] \leftarrow \arg \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

**return** the path by following backpointers from  $backpointer[N, L + 1]$

$L$

$L$

$N$   
 $L$



# Pseudocode for the Viterbi Algorithm



**Input:** observations of length  $N$ , state set of length  $L$

**Output:** best-path

create a path probability matrix  $viterbi[N, L + 2]$

create a path backpointer matrix  $backpointer[N, L + 2]$

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$

$backpointer[1, s] \leftarrow 0$

**end**

**for each** time step  $i$  from 2 to  $N$  **do**

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[i, s] \leftarrow \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s) \times emit(o_i, s)$

$backpointer[i, s] \leftarrow \arg \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s)$

**end**

**end**

$viterbi[N, L + 1] \leftarrow \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

$backpointer[N, L + 1] \leftarrow \arg \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

**return** the path by following backpointers from  $backpointer[N, L + 1]$

$L$

$N$

$L$

$L$

$L$

# Pseudocode for the Viterbi Algorithm



**Input:** observations of length  $N$ , state set of length  $L$

**Output:** best-path

create a path probability matrix  $viterbi[N, L + 2]$

create a path backpointer matrix  $backpointer[N, L + 2]$

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$

$backpointer[1, s] \leftarrow 0$

**end**

**for each** time step  $i$  from 2 to  $N$  **do**

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[i, s] \leftarrow \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s) \times emit(o_i, s)$

$backpointer[i, s] \leftarrow \arg \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s)$

**end**

**end**

$viterbi[N, L + 1] \leftarrow \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

$backpointer[N, L + 1] \leftarrow \arg \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

**return** the path by following backpointers from  $backpointer[N, L + 1]$

$L$

$N$

$L$

$L$

$$L + L^2N$$

# Pseudocode for the Viterbi Algorithm



**Input:** observations of length  $N$ , state set of length  $L$

**Output:** best-path

create a path probability matrix  $viterbi[N, L + 2]$

create a path backpointer matrix  $backpointer[N, L + 2]$

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$

$backpointer[1, s] \leftarrow 0$

**end**

**for each** time step  $i$  from 2 to  $N$  **do**

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[i, s] \leftarrow \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s) \times emit(o_i, s)$

$backpointer[i, s] \leftarrow \arg \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s)$

**end**

**end**

$viterbi[N, L + 1] \leftarrow \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

$backpointer[N, L + 1] \leftarrow \arg \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

**return** the path by following backpointers from  $backpointer[N, L + 1]$

$L$

$N$

$L$

$L$

$N$

$$L + L^2N + N$$

# Pseudocode for the Viterbi Algorithm



**Input:** observations of length  $N$ , state set of length  $L$

**Output:** best-path

create a path probability matrix  $viterbi[N, L + 2]$

create a path backpointer matrix  $backpointer[N, L + 2]$

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$

$backpointer[1, s] \leftarrow 0$

**end**

**for each** time step  $i$  from 2 to  $N$  **do**

**for each** state  $s$  from 1 to  $L$  **do**

$viterbi[i, s] \leftarrow \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s) \times emit(o_i, s)$

$backpointer[i, s] \leftarrow \arg \max_{s'=1}^L viterbi[i-1, s'] \times trans(s', s)$

**end**

**end**

$viterbi[N, L + 1] \leftarrow \max_{s=1}^L viterbi[s, N] \times trans(s, \langle /S \rangle)$

$backpointer[N, L + 1] \leftarrow \arg \max_{s=1}^L viterbi[N, s] \times trans(s, \langle /S \rangle)$

**return** the path by following backpointers from  $backpointer[N, L + 1]$

$L$

$N$

$L$

$L$

$N$

$$O(L^2N)$$

The HMM models the process of generating the labelled sequence. We can use this model for a number of tasks:

- ▶  $P(S, O)$  given  $S$  and  $O$
- ▶  $P(O)$  given  $O$
- ▶  $S$  that maximises  $P(S|O)$  given  $O$
- ▶  $P(s_x|O)$  given  $O$
- ▶ We can also learn the model parameters, given a set of observations.

## Task

Given an observation sequence  $O$ , determine the likelihood  $P(O)$ , according to the HMM.

## Task

Given an observation sequence  $O$ , determine the likelihood  $P(O)$ , according to the HMM.

Compute the **sum over all possible state sequences**:

$$P(O) = \sum_S P(O, S)$$

For example, the ice cream sequence 3 1 3:

$$\begin{aligned} P(3 \ 1 \ 3) = & P(3 \ 1 \ 3, \text{cold cold cold}) + \\ & P(3 \ 1 \ 3, \text{cold cold hot}) + \\ & P(3 \ 1 \ 3, \text{hot hot cold}) + \dots \Rightarrow O(L^N N) \end{aligned}$$



Again, we use **dynamic programming**—storing and reusing the results of partial computations in a **trellis**  $\alpha$ .





Again, we use **dynamic programming**—storing and reusing the results of partial computations in a **trellis**  $\alpha$ .

Each cell in the trellis stores the probability of being in state  $s_x$  after seeing the first  $i$  observations:

$$\begin{aligned}\alpha_i(x) &= P(o_1 \dots o_i, s_i = x) \\ &= \sum_{k=1}^L \alpha_{i-1}(k) \cdot P(x|k) \cdot P(o_i|x)\end{aligned}$$




Again, we use **dynamic programming**—storing and reusing the results of partial computations in a **trellis**  $\alpha$ .

Each cell in the trellis stores the probability of being in state  $s_x$  after seeing the first  $i$  observations:


$$\begin{aligned}\alpha_i(x) &= P(o_1 \dots o_i, s_i = x) \\ &= \sum_{k=1}^L \alpha_{i-1}(k) \cdot P(x|k) \cdot P(o_i|x)\end{aligned}$$

Note  $\sum$ , instead of the max in Viterbi.


# An Example of the Forward Algorithm




# An Example of the Forward Algorithm




# An Example of the Forward Algorithm




# An Example of the Forward Algorithm



# An Example of the Forward Algorithm




# An Example of the Forward Algorithm







# An Example of the Forward Algorithm



# An Example of the Forward Algorithm



# An Example of the Forward Algorithm



$$P(3 \ 1 \ 3) = 0.0033172$$

# Pseudocode for the Forward Algorithm



**Input:** *observations of length  $N$ , state set of length  $L$*

**Output:** *forward-probability*

create a probability matrix  $forward[N, L + 2]$

**for each** *state  $s$  from 1 to  $L$*  **do**

$forward[1, s] \leftarrow trans(\langle S \rangle, s) \times emit(o_1, s)$

**end**

**for each** *time step  $i$  from 2 to  $N$*  **do**

**for each** *state  $s$  from 1 to  $L$*  **do**

$forward[i, s] \leftarrow$   
         $\sum_{s'=1}^L forward[i-1, s'] \times trans(s', s) \times emit(o_i, s)$

**end**

**end**

$forward[N, L + 1] \leftarrow \sum_{s=1}^L forward[N, s] \times trans(s, \langle /S \rangle)$

**return**  $forward[N, L + 1]$

To evaluate a part-of-speech tagger (or any classification system) we:

- ▶ train on a labelled training set
- ▶ test on a *separate* test set

For a POS tagger, the standard evaluation metric is tag accuracy:

$$Acc = \frac{\text{number of correct tags}}{\text{number of words}}$$

The other metric sometimes used is *error rate*:

$$error\ rate = 1 - Acc$$

- Part-of-speech tagging as an example of sequence labelling.

- ▶ Part-of-speech tagging as an example of sequence labelling.
- ▶ Hidden Markov Models to model the observation and hidden sequences.

- ▶ Part-of-speech tagging as an example of sequence labelling.
- ▶ Hidden Markov Models to model the observation and hidden sequences.
- ▶ Learn the parameters of HMM (i.e. transition and emission probabilities) using MLE.





- ▶ Part-of-speech tagging as an example of sequence labelling.
- ▶ Hidden Markov Models to model the observation and hidden sequences.
- ▶ Learn the parameters of HMM (i.e. transition and emission probabilities) using MLE.
- ▶ Use Viterbi for decoding, i.e.:  $S$  that maximises  $P(S|O)$  given  $O$ .



- ▶ Part-of-speech tagging as an example of sequence labelling.
- ▶ Hidden Markov Models to model the observation and hidden sequences.
- ▶ Learn the parameters of HMM (i.e. transition and emission probabilities) using MLE.
- ▶ Use Viterbi for decoding, i.e.:  $S$  that maximises  $P(S|O)$  given  $O$ .
- ▶ Use Forward for computing likelihood, i.e.:  $P(O)$  given  $O$